iCog Labs Artificial Intelligence Technical Report
Transformers and Self-Attention Mechanism.
Vaswani et al. (2017) introduced the Transformer architecture based on self-attention.
Query, Key, and Value vectors project semantic representations.
Redis caching provides cache-aside answer storage to optimize throughput.